Sequence models & Attention mechanism
Quiz, 10 questions

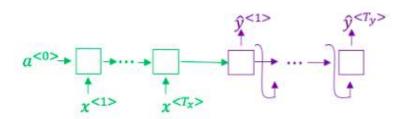✓ **Congratulations! You passed!**

Next Item

✓   1 / 1
point

1.
Consider using this encoder-decoder model for machine translation.

This model is a "conditional language model" in the sense that the encoder portion (shown in green) is modeling the probability of the input sentence x.

○ True

○ False

Correct

✓   1 / 1
point

2.
In beam search, if you increase the beam width B, which of the following would you expect to be true? Check all that apply.

☐ Beam search will run more slowly.

Correct

Sequence models & Attention mechanism

Quiz, 10 questions

☐ Beam search will use up more memory.

Correct

☐ Beam search will generally find better solutions (i.e. do a better job maximizing $P(y \mid x)$)

Correct

☐ Beam search will converge after fewer steps.

Un-selected is correct

✓ 1 / 1
point

3.

In machine translation, if we carry out beam search without using sentence normalization, the algorithm will tend to output overly short translations.

◉ True

Correct

◯ False

✓ 1 / 1
point

4.

Suppose you are building a speech recognition system, which uses an RNN model to map from audio clip x to

a text transcript y. Your algorithm uses beam search to try to find the value of y that maximizes $P(y \mid x)$.

On a dev set example, given an input audio clip, your algorithm outputs the transcript $\hat{y} = $ "I'm building an A Eye system in Silly con Valley.", whereas a human gives a much superior transcript $y^* = $ "I'm building an AI system in Silicon Valley."

According to your model,

$$P(\hat{y} \mid x) = 1.09 * 10^-7$$

$$P(y^* \mid x) = 7.21 * 10^-8$$

Would you expect increasing the beam width B to help correct this example?

○   No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.

▲

Correct

○   No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.

○   Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.

○   Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.

✔   1 / 1 point

5.

Continuing the example from Q4, suppose you work on your algorithm for a few more weeks, and now find that for the vast majority of examples on which your algorithm makes a mistake, $P(y^* \mid x) > P(\hat{y} \mid x)$. This suggest you should focus your attention on improving the search algorithm.
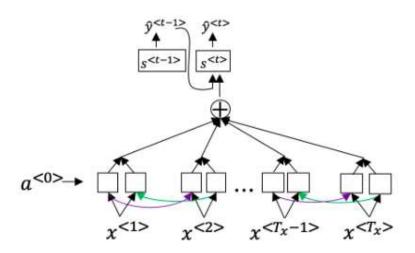
○   True.

▲

Correct

○   False.

Sequence models & Attention mechanism

1 / 1
point

9/10 points (90%)

Quiz, 10 questions

6.

Consider the attention model for machine translation.

Further, here is the formula for $\alpha^{<t,t'>}$.

$$\alpha^{<t,t'>} = \frac{\exp(e^{<t,t'>})}{\sum_{t'=1}^{T_x} \exp(e^{<t,t'>})}$$

Which of the following statements about $\alpha^{<t,t'>}$ are true? Check all that apply.

- [ ] We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t'>}$ that are highly relevant to the value the network should output for $y^{<t>}$. (Note the indices in the superscripts.)

  Correct

- [ ] We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t>}$ that are highly relevant to the value the network should output for $y^{<t'>}$. (Note the indices in the superscripts.)

  Un-selected is correct

- [ ] $\sum_t \alpha^{<t,t'>} = 1$ (Note the summation is over t.)

  Un-selected is correct

☐ $\sum_{t'} \alpha^{<t,t'>} = 1$ (Note the summation is over t'.)

Sequence models & Attention mechanism
Quiz, 10 questions

9/10 points (90%)

Correct

✔   1 / 1
point

7.

The network learns where to "pay attention" by learning the values $e^{<t,t'>}$, which are computed using a small neural network:

We can't replace $s^{<t-1>}$ with $s^{<t>}$ as an input to this neural network. This is because $s^{<t>}$ depends on $\alpha^{<t,t'>}$ which in turn depends on $e^{<t,t'>}$; so at the time we need to evalute this network, we haven't computed $s^{<t>}$ yet.

○   True

▲

Correct

○   False

✔   1 / 1
point

8.

Compared to the encoder-decoder model shown in Question 1 of this quiz (which does not use an attention mechanism), we expect the attention model to have the greatest advantage when:

○   The input sequence length T_x is large.

▲

Correct

○   The input sequence length T_x is small.

✖   0 / 1
point

9.

Under the CTC model, identical repeated characters not separated by the "blank" character (_) are collapsed. Under the CTC model, what does the following string collapse to?

__c_oo_o_kk___b_ooooo__oo__kkk

○ cokbok

○ cookbook

○ cook book

⊘ coookkboooooookkk

This should not be selected

✔   1 / 1
        point

10.
In trigger word detection, $x^{<t>}$ is:

○ Features of the audio (such as spectrogram features) at time t.

Correct

○ The t-th input word, represented as either a one-hot vector or a word embedding.

○ Whether the trigger word is being said at time t.

○ Whether someone has just finished saying the trigger word at time t.